

Nelo Operations Exercise

Thank you for taking the time to complete this exercise! The prompt below is intentionally open-ended, so please ask clarifying questions if needed. Feel free to organize your response in a document or deck, whichever you feel will best frame your thinking. We look forward to reviewing your work together!

We expect you to use AI tools throughout, including AI coding tools (Claude Code, Codex, Cursor, or similar) for the build portion. This is a real part of how the role works.

Hot Sale Readiness

Prompt:

You are joining Nelo 3 weeks before Hot Sale. Marketplace growth is accelerating, but operational performance is deteriorating: SLA compliance has dropped from 92% to 71%, the support backlog is 4x what it was 8 weeks ago, and ~60% of the ops team's time goes to manual reconciliation and report-pulling. Each lead tracks their own metrics in their own sheet. You cannot add headcount before Hot Sale, but you have AI tools and ~10% of engineering bandwidth (Attached is a sample dataset, ~90 days of ops activity).

Important context:

- GMV is expected to be 3–4x baseline during Hot Sale week
- The CEO wants you to use this as a forcing function to professionalize how ops runs, not just survive the event
- AI tooling, including AI coding tools, is expected to be a core part of how this team operates

Ask:

1. **Diagnose.** Use the data to ground your reading of what's happening. What's the real story? What are you fixing in the next 3 weeks, what are you deferring, and what does "good enough" look like for Hot Sale itself?
2. **Measure.** Define the metrics and reporting structure ops will run on after Hot Sale, what gets measured, by whom, at what cadence, and how decisions actually get made from it.
3. **Build.** Using Claude Code, Codex, or similar, build one thing that would meaningfully change how this ops team works (a dashboard), a reconciliation tool, an alerting workflow, a triage classifier, whatever you'd ship first. Share the working artifact (repo or runnable code) and the AI conversations that built it. Be ready to walk us through how you decomposed the problem, how you directed the model, where you caught it being wrong, and what you'd build next.

ASSUMPTIONS:

The dataset contains no SLA policy, so every breach figure in this document rests on an explicit assumption: SLA = time to first response, with thresholds of 4 hours (high priority), 8 hours (medium), and 24 hours (low). Resolution-time SLA cannot be measured because the dataset has no resolved_at field. All breach percentages are sensitive to this assumption, which is why the triage tool exposes these thresholds as editable settings — if Nelo's actual policy differs, the analysis recalculates rather than breaks.

Diagnose:

Marketplace growth exposed a payments operations bottleneck. The support crisis is being driven primarily by **payment operations** rather than customer support itself. Representatives are buried in manual spreadsheets trying to resolve systemic transactional anomalies. The current process relies heavily on manual reconciliation and fragmented reporting, causing backlog growth and SLA deterioration.

Primary Root Causes

- **Reconciliation Infrastructure is broken and the current reconciliation processes are breaking under volume.**
 - Ticket volume grew ~80% from Dec → Feb.
 - Reconciliation exceptions x4 from Dec → Feb
 - Pending transactions spiked dramatically (x16 from Dec → Feb)
 - 3-day incident (Feb 9–11) generated ~60% of all pending transactions
- **OXXO & Ledger Logic Faults**
 - OXXO consistently shows the highest reconciliation exception behavior (48% of all exceptions). This is likely generating a large share of payment tickets, but due to fragmented reporting is difficult to know.
 - Internal ledger is generating inverted negative signs for ~1% of all transactions (e.g., expected negative values matched against positive settled amounts). This causes skewed merchant balances, resulting in high-urgency billing tickets.
- **Lack of operational visibility**
 - Since each lead tracks their own metrics in their own sheet, there is no unified view of end-to-end resolution times or early warning system.
- **Poor prioritization**
 - Tickets from top merchants are taking the same average time as small merchants.
 - All “issue types” have the same urgency/priority.
 - Nothing escalates with time: a ticket about to breach (or already past) its SLA window looks identical in the queue to one that just arrived.

The 3-Week Pre-Hot Sale Action Plan

Fix #1: AI Automated Reconciliation Pipeline (Week 1 — Owner: Payments Lead + engineering)

Deploy a Python script via engineered workflows to pull daily transaction files. It will programmatically resolve inverted sign flags, isolate genuine mismatches, and auto-enrich records with merchant database names.

Fix #2: Clear Backlog with AI Triage Bots (Week 2 — Owner: Support Lead)

Deploy AI bots to instantly handle high-frequency, low-risk requests. For example, for “*account_access*” bots can instantly cross-reference user data and send password reset/verification links and for “*delivery_delay*” tickets, the bot can auto-ping logistics webhooks and provide status updates. Security guardrail: for *account_access* the bot drafts the response and verifies identity, but a human approves before any reset link is sent — account-access automation is security-sensitive and should not run fully autonomous three weeks before peak volume.

Fix #3: Set up a simple high-level Central Control Dashboard (Week 2–3 — Owner: Ops Lead)

Consolidate database tracking into a single unified dashboard to serve as the team’s operational command center. The dashboard ships as a second view inside the triage tool and doubles as the Hot Sale detection layer: it carries the three threshold tiles (OXXO exception rate, pending rate, payment-issue volume vs. baseline) that the war-room cadence monitors during the event (see the “Good Enough” metrics below).

Fix #4: Build triage tool to prioritize the most urgent tickets (Week 1 — Owner: Ops Lead; built — see Build section)

Re-prioritize tickets assigning different weights to each section

- Category type (*payment_issue*, *refund_request*, *delivery_delay*, *account_access*, other)
- Merchant GMV
- SLA ageing (As tickets spend more time in the queue their priority increases)
- Ticket priority (high, medium, low)

What to defer - post-hot sale backlog

- **Structural Payment Engine Redesigns:** Tweaking core transaction modules three weeks before an event with an expected x4 GMV increase is an unacceptable stability risk.
- **Defer Complex Ticket Auto-Resolutions:** High-value billing/payout disputes must remain manual for safety; do not try to build an AI that moves money right before Hot Sale.
- **Exhaustive dashboard (more details in the next section)**
- **Investigate why Beauty Merchants are generating a higher number of tickets than other marketplace categories**

Hot Sale "Good Enough" Metrics

- **SLA Stability at 80%:** A blended 92% SLA is not achievable at 3–4× volume with fixed headcount — good enough means deciding who waits. Tiered floors: P0 tickets (payment issues + top-GMV merchants) ≥90% first response within 4 hours. P1 ≥80% within 8 hours. P2/P3 explicitly relaxed to 48–72 hours and communicated to merchants in advance rather than silently missed. The blended rate will read ~78–82% and that is accepted.

- **Positive Queue Velocity:** Backlog may grow on peak days but must not compound: day-over-day growth $\leq 25\%$, and shrinking again within 48 hrs after peak. Zero P0 tickets unassigned for more than 30 minutes; P0 queue depth never exceeds 60 (Hot Sale threshold, vs. 30 in steady state).
- **Zero Blind Spots:** Every known failure mode is watched, even if not yet automated. The Central Control Dashboard (Fix #3) carries three threshold tiles
 - 1 OXXO exception rate $>10\%$
 - 2 daily pending rate $>15\%$ (the Feb 9–11 failure mode that ran 3 days undetected and generated $\sim 60\%$ of all pending cases)
 - 3 payment-issue ticket volume $>1.5\times$ baseline

An on-duty lead checks them every 2 hours during Hot Sale week. Building an automated alerting system is deferred to post-Hot Sale (it is the first item in the build-next list).

- **Merchant Communication:** No merchant with >5 open tickets goes 24 hrs without proactive outreach and the top 20 GMV merchants have a named owner for Hot Sale week.

Measure

After Hot Sale, Ops should run through a four-tier centralized Control Ops Dashboard

Tier 1: Real-time Monitoring & Automated Alerts

Objective: Identify if anything is breaking/broken.

Owner: None. No human required to generate – only to act.

Cadence: Metrics are updated every 30 minutes.

KPIs & Actions:

Metric	Threshold	Action
Ticket volume # tickets/hour vs 7-day rolling average for the same hour	> 1.5x of 7-day rolling average for the specific hour	Alert Ops Lead
Payment issue volume # tickets/hour vs 7-day rolling average for the same hour	> 1.5x of 7-day rolling average for the specific hour	Alert Payments Lead
OXXO exception rate rolling-4hr exception rate vs 7-day baseline	Rolling 4-hour exception rate exceeding 15%	Alert Engineering team + Ops lead team
P0 queue depth # of open P0 Tickets	> 50 open P0 tickets in the unassigned pool	Reassign all available agents to P0 queue
Unassigned P0 tickets age	Any P0 ticket left unassigned for >1 hour	Alert team group / channel with merchant name and GMV

Tier 2: Daily Operational Review

Objective: Evaluate the previous day's execution, manage backlog velocity, and balance resource distribution. This becomes the single source of truth but with dedicated sections for each team.

Owner: Ops lead.

Cadence: Daily meeting at the beginning of the day

Metric:

Metric	Threshold	Action
SLA Compliance Rate % of tickets with first response threshold	If SLA compliance rate is <80% for 2 days in a row	Escalate to weekly review + route cause analysis.
Open Backlog volatility Open tickets today vs avg 7 last days	If Backlog growing for 3 days in a row	Triage-weight review + temporary P3 freeze (agents stop pulling P3 tickets) until backlog shrinks for 2 consecutive days
Agent utilization # of open tickets per agent	If any agent's open-ticket count deviates >30% from the team median	Catch load imbalances before one agent becomes the bottleneck
Individual Resolution Velocity Daily closure per agent	If any agent -20% tickets closed vs the average of the group	Review tickets with Reps to see what is going on and how to proceed
Merchants at Risk # of merchants with +3 open tickets	If merchants with >5 open tickets simultaneously	Flags the accounts at risk and triggers a proactive merchant outreach call from the reps.

Tier 3: Weekly Engineering alignment

Objective: Review recurring defects, and prioritize engineering bandwidth

Owner: Ops lead

Participants: Head of Ops, Support Lead, Payments Lead , Merchant Ops Lead, Engineering Representative

Focus/Metrics:

- Section 1: What broke?
- Section 2: Why tickets are being created?
- Section 3: What can be automated?
- Section 4: Updates from Engineering team

Section 1: What broke?

Example:

Problem	Impact
OXXO matching issue	23% exception rate
Beauty merchants	40% of payment tickets
Merchant payout delays	15% SLA misses

Section 2: Root Cause Analysis: why were those tickets created?

Example:

Root Cause	Ticket Share
Settlement mismatch	35%
Payout delay	25%
Merchant misunderstanding	20%
Refund issue	10%
Other	10%

Section 3: Automation Pipeline

Example

Task	Hours Spent	Automation Potential
Reconciliation	25 hrs/week	High
Report Pulling	15 hrs/week	High
Ticket Tagging	10 hrs/week	High

- If a root cause generates >10% of tickets, then create an automation or product fix.
- If a merchant segment generates disproportionate operational cost, adjust process, tooling, pricing, or merchant onboarding.

Section 4: Engineering Updates

Initiative	Status
Triage Tool	Testing – Expecting to be delivered next week
Report Pulling	Not started yet – waiting on API connection
Ticket Auto Tagging	Implemented

Tier 4: Weekly Performance Review

Objective: Review general performance and connect operational performance to unit economics.

Owner: Heads of Ops

Operational Health KPIs:

- **SLA Compliance:** target >90% overall and >85% on payment issues; below 85% for 2 consecutive weeks triggers a category-level root-cause review
- **Average Resolution Time:** how long a merchant waited before their ticket was resolved. Target: declining quarter over quarter; a sustained rise triggers a staffing or automation review.
- **Touches Per Ticket:** # of times a support agent interacts with a single ticket before resolving it. Target: <3; a rising trend points to a process or knowledge-base gap.

- **Avg SLA Breach Depth:** by how many hours the SLA was missed. Target: <2 hours average depth; deep breaches signal abandoned tickets, not slow ones.
- **Re-open rate:** percentage of resolved tickets that merchants reactivate because the initial solution didn't fix their issue. Target: <8%; above that triggers a resolution-quality review of the affected category.
- **Customer Satisfaction Score (CSAT):** Transactional post-resolution feedback surveys. Target: $\geq 4.2/5$; any score below 3 from a top-50 GMV merchant triggers same-week outreach.
- **Tickets per FTE:** monitored as a trend; rising throughput without rising breach rate validates automation gains.
- **Cost per Ticket:** rising for 2 consecutive quarters reprioritizes the automation pipeline.

Build

Since there are 3 weeks before Hot Sale I would focus first on:

- Building triage tool to prioritize the most urgent tickets
- High-level Ops control Dashboard – to identify most pressing issues
- Quick wins – there are some tools that probably are not very time consuming to build and that could save reps some time. For example: account_access issues and delivery_delay issues.

What I built — the triage tool. At a glance, it:

1. Detects reconciliation/payment issues daily
2. Prioritizes the merchants/cases that matter most
3. Creates an actionable queue for Ops, not just charts
4. Groups tickets by merchant — the detail panel shows each merchant's other open tickets and financial exposure

Why this is the right first build:

Right now, a payment issue ticket from MER1058 (Beauty, 626K MXN/month GMV, 36 open tickets) sits in the same queue and gets the same attention as a low-priority "other" ticket from a small merchant. Agents aren't triaging by business impact — they're triaging by arrival order. At 4× volume during Hot Sale, that approach collapses completely.

The classifier manages the symptom efficiently, but it doesn't fix whatever broke. This buys us the capacity to survive Hot Sale while we diagnose the root cause and fix the underlying issue.

Triage Classifier - Auto-tag & prioritize incoming tickets using category + merchant GMV + SLA age

Why the triage classifier first. Four candidates were evaluated — dashboard, reconciliation tool, alerting workflow, triage classifier. Triage won on four criteria:

1. directly attacks the failure mode in the data — payment issues drowning in a flat first-in-first-out queue;
2. scales with volume — at 3–4× tickets, automated prioritization becomes more valuable, while a dashboard only shows the fire in higher resolution;
3. requires zero engineering bandwidth — it connects directly to the team's existing Google Sheet with configurable column mapping and a 5-minute auto-refresh (no export step, no pipeline);
4. protects GMV explicitly.

The reconciliation tool was deferred for the same reason stated in the Diagnose section: shipping financial tooling three weeks before a 4× event is a stability risk. The dashboard was deferred because visibility without changed agent behavior does not move SLA.

How the scoring engine works:

Each ticket gets a 0–100 score from four components:

- priority (0–30);

- category risk (0–30 — weights derived from observed breach rates, payment_issue gets the maximum because it breaches at ~31% vs. 4–8% for all other categories);
- merchant GMV (0–25, scored in the same size bands used in the diagnosis — >750K=25, 500–750K=20, 250–500K=15, 100–250K=8, <100K=3 — capped so a single large merchant cannot bury the rest of the queue);
- SLA aging (0–15, rising as a ticket approaches and passes its threshold so nothing rots silently).

Tiers: P0 ≥75 (escalate now), P1 ≥55, P2 ≥35, P3 below. Every weight and threshold is shown in “Scoring engine settings” panel, because the weights are calibrated assumptions rather than ground truth — ops can recalibrate without touching code.

Beyond ranking tickets — merchant financial exposure

The tool also merges the Reconciliation information into the queue. Each merchant receives a financial-exposure tier (CRITICAL / HIGH / MEDIUM) calculated from settlement difference, exception count, active payment tickets, and merchant size; the ticket detail panel shows the merchant’s other open tickets and their exposure; and the dashboard tracks total cash at risk. This is the artifact-level connection between the two problems in the diagnosis — ticket SLA collapse and reconciliation exceptions — in a single view.

How I decomposed the problem and directed the AI

In order: (1) explored the dataset to understand what was happening before building anything; (2) explored potential reasons for the issues and the impact of different solution to "fix it " or make it through the Hot-Sale. I ran the same problem through ChatGPT, Claude, and Gemini, compared the pros and cons of each proposal, and landed on a hybrid structure combining the strongest elements of each; (3) built the UI of a working ranked queue with a per-ticket score breakdown, so agents can see why a ticket ranks where it does; (4) added a settings layer so weights and thresholds can be changed without touching code — if something else breaks during Hot Sale week, the triage tool can be recalibrated on the spot; (5) added the dashboard view for the high-level picture, including the three Hot Sale watch tiles, without going deep on any single ticket.

Where I caught the AI being wrong

- Initial date parsing produced an impossible timeline from Excel serial dates — caught by checking min/max dates against the stated 90-day window; the dataset was corrected.
- The model presented SLA breach percentages as fact when the underlying thresholds were assumptions it had invented — I challenged the definition.
- The model claimed Beauty merchants over-index on OXXO to explain the Beauty ticket concentration — my own cross-tab showed payment-method mix is identical (~24% OXXO) across all categories, so the explanation did not hold; the real finding is a small cluster of high-volume Beauty merchants, which is why the open investigative question below is to read their actual ticket text.

What I would build next. A volume/pending anomaly alert tool that automates the Hot Sale war-room checks — a script that would have caught both the Feb 9–11 incident and the OXXO drift in hour one, and that becomes the engine behind Tier 1 of the post-Hot Sale Measure framework.

APPENDIX

Diagnose – Analysis

Ticket Volume and Avg response time

Ticket volume increasing rapidly, +80% from December to February but the biggest issue is that the Average response time is getting bigger in February (x4).

Month	Tickets (#)	Avg response time
Dec-24	1,228	2.83
Jan-25	1,922	3.11
Feb-25	2,214	12.26

Ticket volume doesn't seem to be growing rapidly but in January 31st we start to see an increase in the Avg. Response time

Week Number	Tickets (#)	Avg response Time	Tickets Growth WoW (#)	Tickets Growth WoW (%)
2024 - W49	238	2.98		
2024 - W50	232	2.72	-6	-3%
2024 - W51	291	2.66	59	25%
2024 - W52	315	2.95	24	8%
2025 - W01	358	2.94	43	14%
2025 - W02	365	2.54	7	2%
2025 - W03	446	2.70	81	22%
2025 - W04	478	2.90	32	7%
2025 - W05	516	5.95	38	8%
2025 - W06	513	11.15	-3	-1%
2025 - W07	525	12.56	12	2%
2025 - W08	629	12.23	104	20%
2025 - W09	458	12.90	-171	-27%

This signals potential issues:

- Staffing may not have scaled with ticket volume
- Backlogs may be forming

Employee performance

No obvious employee performance problem. The five agents (Alice, Bob, Carol, David, Eve) have very similar average response times and ticket assigned to them, even after January 31st.

Employee	Tickets (#)	Avg Response time
Carol	1,008	6.5
Bob	993	6.7
Alice	992	6.8
Eve	948	7.0
David	1,010	7.2

Ticket priority

High-priority tickets are not being responded to materially faster than low-priority tickets. **Business should expect priority based service levels.**

Priority	Avg response time
High	7.1
Medium	6.9
Low	6.7

Merchant size & type

The issue does not seem to be linked to a particular Merchant size or category.

Avg Monthly GMV (MXN)	Tickets (#)	Merchants (#)	Tickets per merchant	Avg response time
>750K	148	6	25	7.84
500 - 750K	655	12	55	8.00
250-500K	647	12	54	6.64
100-250K	2,642	82	32	6.82
<100K	1,252	38	33	6.15

Issue: the biggest merchants cannot have the same average response time as a small merchant.

Beauty merchants generate disproportionate ticket volume. This suggests the pain is concentrated in a **small group of high-volume Beauty merchants**, not evenly distributed across the marketplace.

Merchant Category	Tickets (#)	Merchants (#)	Tickets per merchant	Avg response time
Beauty	2,223	16	139	7.07
Electronics	1,151	50	23	6.34
Fashion	941	41	23	6.01
Home	710	30	24	7.43
Other	319	13	25	7.68

Category

When looking at the category we notice that the issue comes from ticket categorized as “payment_issue”

Category	Tickets (#)	Avg response Time	% Over 24 hours
payment_issue	1,537	14.3	20%
delivery_delay	1,279	4.0	0%
refund_request	1,051	4.0	0%
account_access	708	3.9	0%
other	789	3.3	0%

In February, payment issue response time explodes;

Month	Payment tickets (#)	Avg response time	% over 8 hrs	% over 24 hrs
Dec-24	269	3.0 hrs	6.7%	0.0%
Jan-25	415	4.0 hrs	6.5%	2.7%
Feb-25	853	22.8 hrs	67.6%	35.6%

Marketplace Category	Payment Issue Tickets	% of Payment Issue Tickets	Avg Response Time
Home	208	14%	16.5 hrs
Other	98	6%	16.2 hrs
Beauty	633	41%	15.2 hrs
Electronics	335	22%	12.7 hrs
Fashion	259	17%	11.8 hrs

Beauty is not the slowest category, but it contributes the largest absolute number of problematic tickets because of its volume.

Most tickets come from merchants that have been live for **1+ year**, likely because they have more volume and more payment activity.

Merchant age Tickets Avg response

0–3 months 60 3.1 hrs

3–6 months 624 5.1 hrs

6–12 months 1,014 7.6 hrs

1–2 years 3,209 6.5 hrs

2+ years 437 10.6 hrs

Beauty is not generating a higher percentage of Payment issues.

Row Labels	account_access	delivery_delay	other	payment_issue	refund_request
Beauty	14%	23%	14%	28%	20%
Electronics	13%	25%	14%	29%	19%
Fashion	13%	24%	16%	28%	20%
Home	12%	24%	16%	29%	18%
Other	15%	22%	12%	31%	20%

Reconciliation

Reconciliation transactions increase (+, but exception rate increases much faster:

Month	Transactions	Exception rate	Pending rate
December	3,281	3.7%	1.0%
January	5,015	4.6%	2.2%
February	5,743	10.4%	9.7%

On Feb 9–11, reconciliation breaks badly

Date	Tx	Exception rate	Pending rate
01-Feb-25	205	12%	1%
02-Feb-25	183	8%	3%
03-Feb-25	204	9%	1%
04-Feb-25	181	8%	3%
05-Feb-25	177	11%	2%
06-Feb-25	233	12%	3%
07-Feb-25	196	12%	3%
08-Feb-25	167	10%	4%
09-Feb-25	196	12%	70%

10-Feb-25	217	11%	69%
11-Feb-25	182	9%	73%
12-Feb-25	187	9%	3%
13-Feb-25	220	10%	3%
14-Feb-25	197	9%	0%
15-Feb-25	238	12%	3%

Those 3 days alone explain **419 pending transactions**, about **60% of all pending cases**. This looks like a **system/process incident**.

Payment method

OXXO is the biggest payment-method outlier with the higher exception rate.

Payment Method	Transactions	Exception rate	Pending rate
OXXO	3,413	13%	6%
Card	6,387	4%	5%
SPEI	2,830	5%	5%
Wallet	1,409	5%	4%

In February, OXXO # of exceptions growth faster than the rest of payment methods, explaining the peak in February.

Week Number	Card	OXXO	SPEI	Wallet	Total exception
2024 - W49	6	7	5	2	20
2024 - W50	15	5	2	1	23
2024 - W51	11	7	5	2	25
2024 - W52	13	9	6	6	34
2025 - W01	18	16	12	8	54
2025 - W02	21	13	8	6	48
2025 - W03	17	11	10	5	43
2025 - W04	24	9	11	3	47
2025 - W05	33	32	12	5	82
2025 - W06	32	86	11	5	134
2025 - W07	37	77	21	15	150
2025 - W08	35	111	14	13	173
2025 - W09	24	74	11	6	115
Total	286	457	128	77	948

OXXO is not to blame for the Feb 9-11 issue, since the impact was across all payment methods.

Date	Card	OXXO	SPEI	Wallet	Total pending tx
08-Feb-25	3	3		1	7
09-Feb-25	74	29	25	9	137
10-Feb-25	64	48	20	17	149
11-Feb-25	50	37	32	14	133
12-Feb-25	2		2	1	5

There are likely two problems:

1. A broad Feb 9–11 reconciliation incident.
2. A persistent OXXO-specific reconciliation issue.

The exception rate or pending rate is not linked to any particular merchant category

Merchant category	Transactions	Exception rate	Pending rate
Beauty	1,536	6%	5%
Electronics	4,669	7%	5%
Fashion	3,839	6%	5%
Home	2,787	7%	5%
Other	1,208	7%	5%

SLA Breach status category	SLA Breach (Multiple Items) (All)
----------------------------	-----------------------------------

merchant_id	Marketplace Category	Sum of first_response_hours
<u>MER1141</u>	Beauty	383.71
<u>MER1005</u>	Beauty	294.73
<u>MER1122</u>	Beauty	260.57
<u>MER1047</u>	Beauty	235.92
<u>MER1129</u>	Beauty	219.5
<u>MER1046</u>	Beauty	180.02

MER1058	Beauty	143.59
MER1120	Beauty	129.9
MER1108	Home	114.98
MER1044	Electronics	89.12

8 out of 10 merchant in the backlog (how do we define backlog) are from the Beauty Category

This suggests the root cause is tied to a specific integration, payment processor, or cohort of merchants rather than a platform-wide issue integration that Beauty merchants use disproportionately.

Beauty is not having more OXXO payments than other categories

	card	oxxo	spei	wallet
Beauty	45%	24%	20%	11%
Electronics	46%	24%	20%	10%
Fashion	46%	24%	20%	10%
Home	44%	25%	20%	10%
Other	44%	25%	20%	11%

Not an onboarding issue. Merchants in their first 3 months generated only 60 tickets out of 5,364 total (~1%). The majority of tickets come from merchants that have been on the platform more than 1 year. So even if a new merchant is less familiar with the platform, they often create fewer tickets simply because they have much less activity.

- The problem is concentrated among a small cluster of high-volume Beauty merchants whose payment-related tickets are taking 15–21 hours to receive a first response.
- # of volume growing since December but Avg response time shoots in February 9,10 and 11

Open investigative questions (deliberate next step, not an oversight)

- Pull all tickets for MER1141, MER1046, MER1047, MER1058, MER1129, and MER1005. Read the ticket text. Determine whether they are reporting the same underlying payment issue (settlement delays, failed payouts, reconciliation problems, chargebacks, etc.).